Deploy a Model Using JumpStart

To Begin with the Lab

Summary of the Lab

Open Amazon SageMaker Studio, create and run a private JupyterLab space named test-space, then access JumpStart Foundation Models. Select and deploy an image classification model using ml.m5.xlarge. Test inference, open the generated notebook in JupyterLab, and execute the pre-written inference code successfully.

- Open Amazon SageMaker Studio.
- Navigate to JupyterLab.
- Click Create Space.

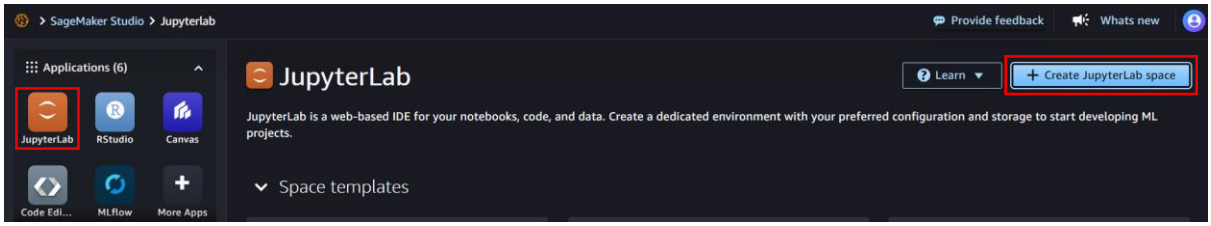

- Enter:
- Name: test-space
- Type: Private
- Click Create
- Start the space and wait until status is Running

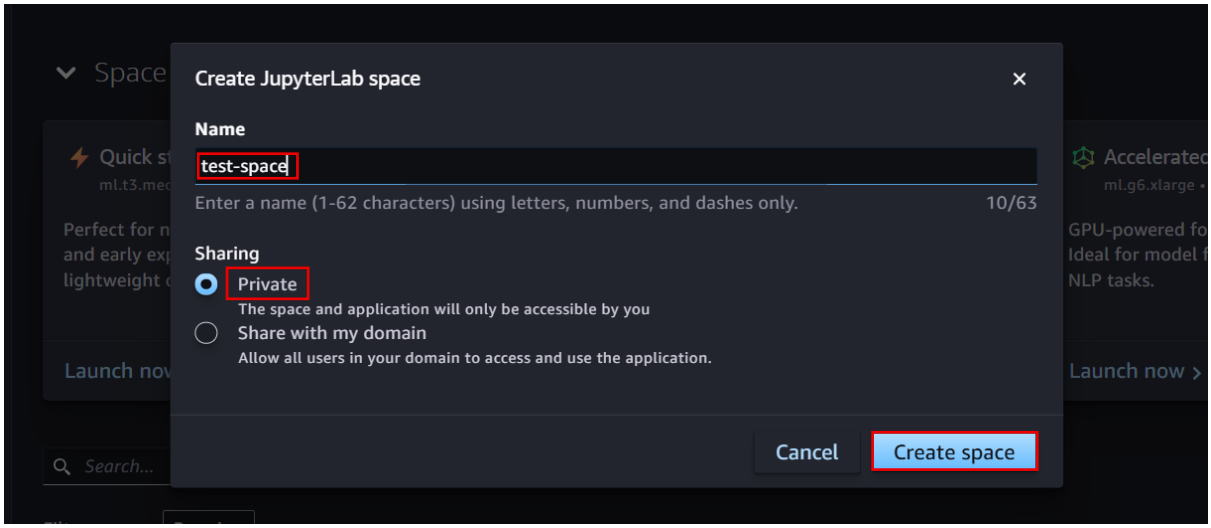

- Click Run Space

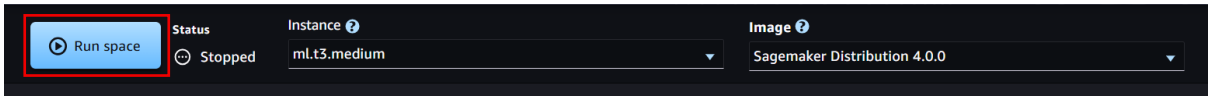

- Click Open JupyterLab once space is ready.

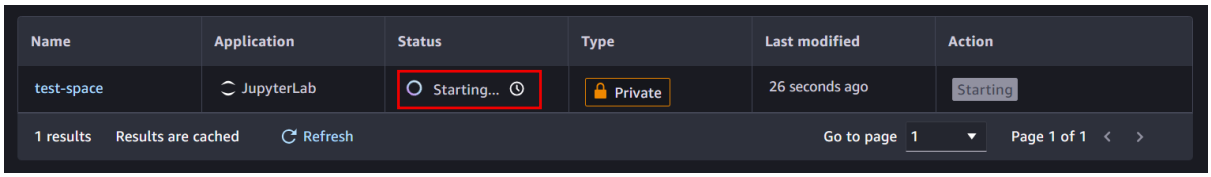

- Click open

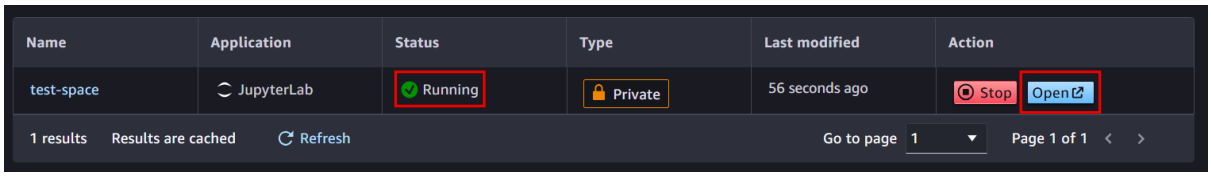

- Open JumpStart
- In SageMaker Studio, go to:
- JumpStart → Foundation Models
- Browse or search:
- Example: image classification

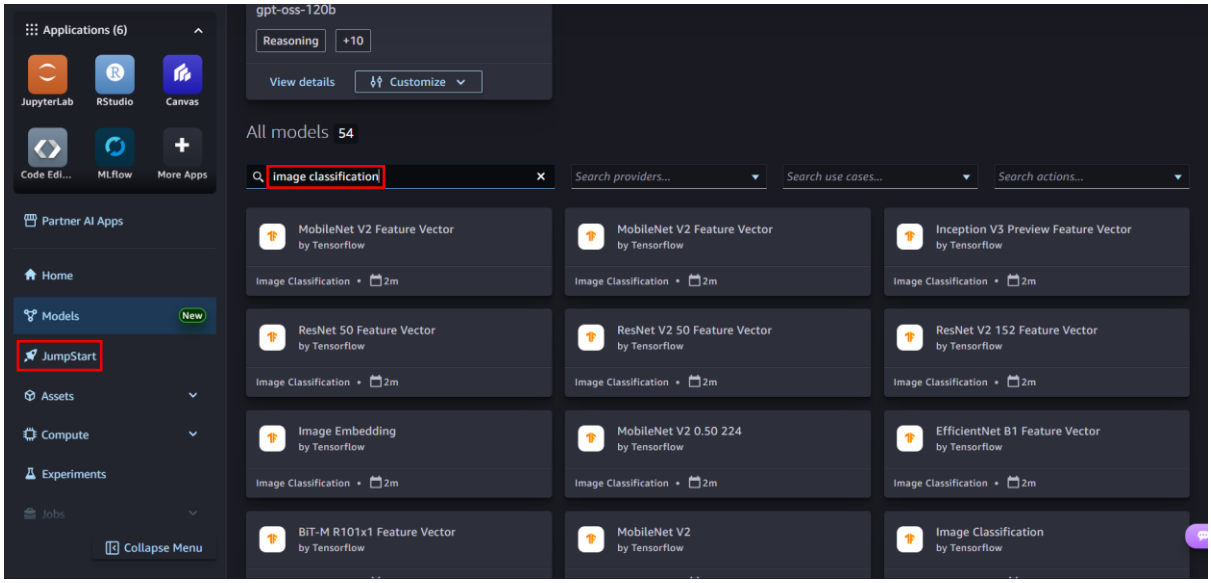

- Choose an image classification model
- Click Deploy

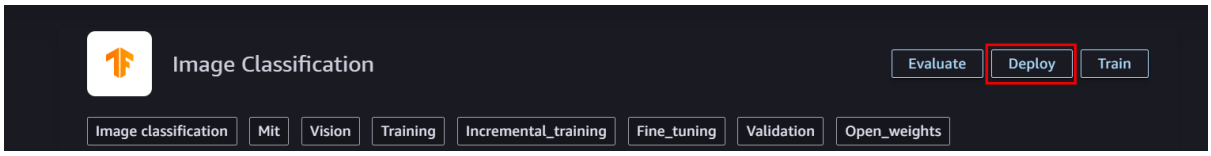

- Configure:
- Endpoint name (default is fine)
- Instance type: ml.m5.xlarge (free tier eligible)
- Instance count: 1

- Click Deploy

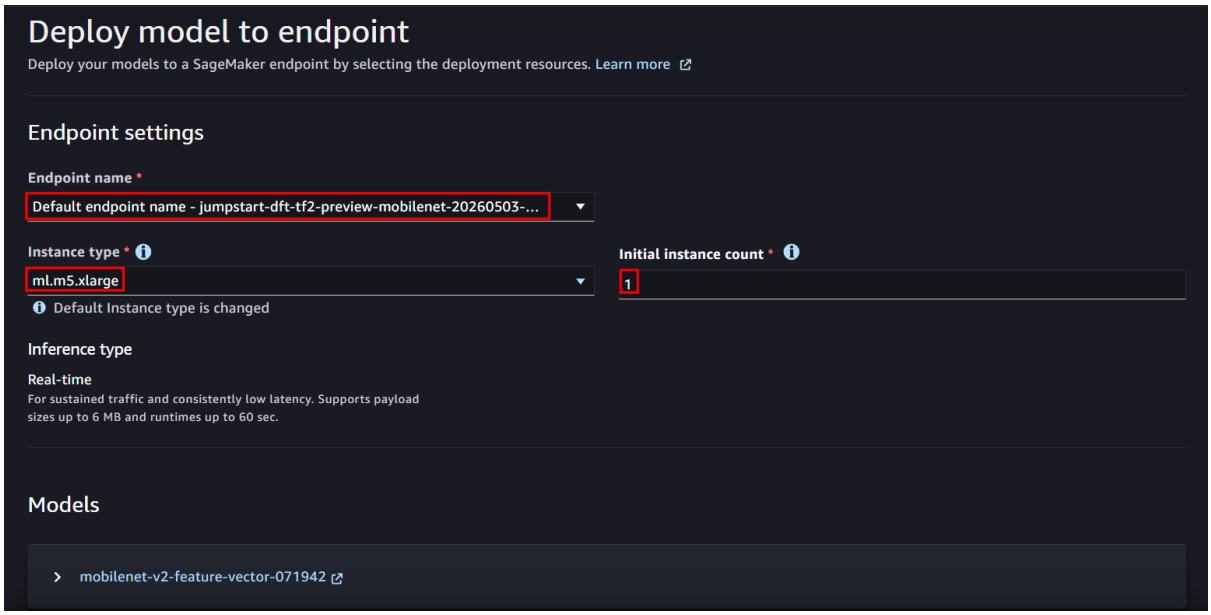

- Click Test Inference
- Open in your running JupyterLab space
- Notebook opens with pre-written inference code
- And run the code Jupyter notebook under Jupyter Lab.